

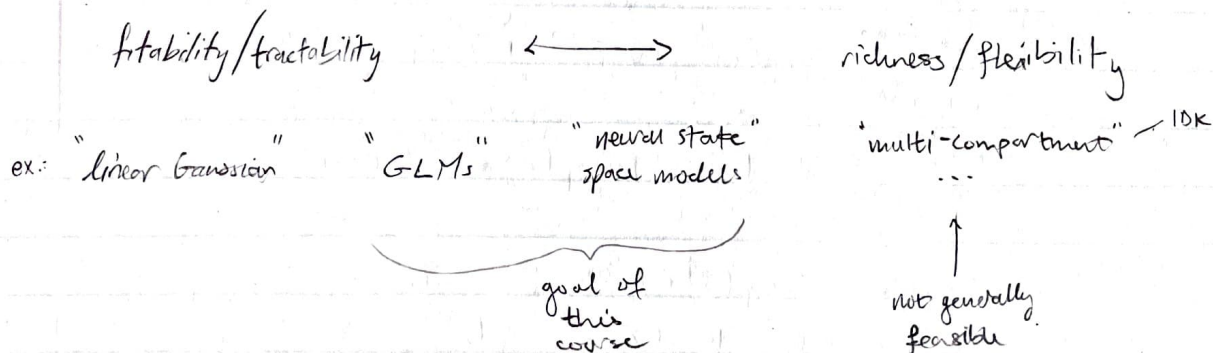
Intro & Course Background

- "modular" structure of probabilistic models
 - ↳ plug & play pieces
- framework for thinking about prob. modeling
- this course is just a starting point → introduce the main components
- TODO b/w sessions:
 1. reading book & paper(s) ← make sure to use "Perusal" webapp.
 2. exercises & "brief write-up"
 3. complete quiz by noon of day of class (diagnostic)

Modeling

- build a model for $p(y|x)$
 - ↳ "statistical encoding model"
- latent variable model (lower dimensional latent representation)
 - w/ unknown variable z believed to exist: $p(y|z)p(z)$
 - can extend to time: $p(y_t|z_t)p(z_t|z_{t-1})$ "latent dynamics"

Model desiderata



Bayesian recap


1. prob. & prob. models
2. inference of prob. models
3. model comparison

1. Probability & Probabilistic Models

- stimulus x & response y
- want a parametric distribution to describe $p(y|x)$

	$p(y x)$	parameters
Poisson	$\text{Poisson}(y \lambda(x))$	$\lambda(x)$
Gaussian	$N(y \mu(x), \sigma^2(x))$	$\mu(x), \sigma^2(x)$

- prob. dist: a fcn that maps from x to a probability
 - $p(X=x) \equiv p(x)$ is the prob. that r.v. X is x
 - for a PDF, $p(x) \in [0, \infty]$
 - for a PMF, $p(x) \in [0, 1]$
- transforming b/w two r.v.:



$$p(x)dx = p(y)dy$$

$$p(y) = p(x) \left| \frac{dx}{dy} \right|$$

$$p(y) = p(x) \left| \frac{1}{f'(x)} \right|$$

Probabilities

marginial across y

	x	
	0	1
y	0	0.1 0.2
	1	0.3 0.4

eg. $p(x) = p(x|y=0) + p(x|y=1)$

eg. $p(y|x=1) = \frac{\text{joint}}{p(x=1)}$ (marginal)

independence: ex: $p(y|x) = p(y)$ then y is indep. of x
 x & y are independent if: $p(x, y) = p(x)p(y)$ ← "product rule"

- Bayes rule:
$$\frac{p(y|x) \overset{\text{likelihood}}{p(x)} \overset{\text{prior}}{p(y)}}{p(y)}$$

 (normalize to $\Sigma \rightarrow 1$; marginal likelihood)
 ↳ basic consequence of standard rules of prob.

- Independent: Identically distributed data (i.i.d data)
 · one data point doesn't affect another
 · data come from the same generative dist.

2. Inference w/ probabilistic models

- max. likelihood estimate (MLE)

- find parameters that make the observed data most likely
- maximizes the likelihood

- pros:

- consistent: converge
- efficient to compute

cons:

- noisy for little data
- no estimate of uncertainty

- Bayesian Inference

· most likely model params: $p(\overset{\theta}{\text{params.}} | \overset{y}{\text{data}})$

$$p(\theta | y) = \frac{p(y|\theta) p(\theta)}{p(y)} \propto p(y|\theta) p(\theta)$$

$p(\theta)$: prior belief of prob. of params θ

$p(y|\theta)$: likelihood of data y given params θ

$p(\theta|y)$: posterior prob. of params θ given the data observed y

- maximum a-posteriori inference (MAP)

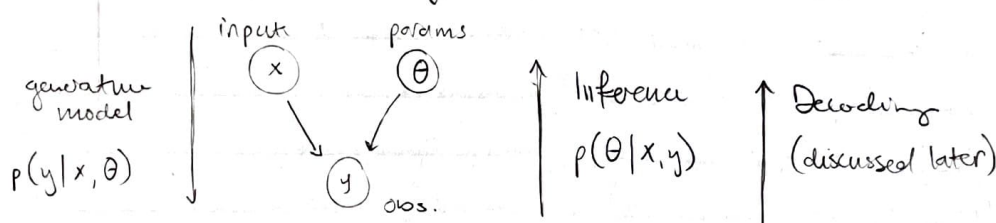
- similar to MLE except maximizing $p(\theta|y) \propto p(\theta)p(y|\theta)$
- same as MLE w/ $p(\theta) \propto 1$

- pros:

- prior helps regularize
- can include prior info

↳ * This is an improper prior & is usually not accurate, anyways

- inference summary



• methods of inference: Full Bayesian, MAP, MLE

3. Model comparison

- Bayesian decision theory

- posterior probs \rightarrow decision

$L(\hat{\theta}, \theta)$: loss function; cost of choosing $\hat{\theta}$ if θ is the real value

$$\hat{\theta}_L = \underset{\hat{\theta}}{\operatorname{argmin}} \int \underbrace{L(\hat{\theta}, \theta)}_{\text{"cost"}} \underbrace{p(\theta|x, y)}_{\text{"weight"}} d\theta$$

$$= \underset{\hat{\theta}}{\operatorname{argmin}} E[L(\hat{\theta}, \theta) | x, y]$$